

Define the data once. Run everything on it.

A Datapipdd paper on method · October 2026

Summary

Most organisations now want AI to do real work: answer clients, run reconciliations, draft reports, move data from an old estate to a new one. Most of those efforts stall. The model is rarely the reason. It stalls because it lands on data nobody defined and on processes nobody wrote down, and because nobody can show afterwards what changed and why.

Datapipdd is an independent data practice. We build the layer underneath: the governed data, the technology that serves it, and the operating model that keeps both true. Every statement our AI work produces keeps the source it came from, passes a separate check, and is cleared by a person before anything moves. This paper describes that method at a high level.

Why AI stalls

The pattern is the same in banks, market infrastructures, retailers and product companies.

- The same customer, product or counterparty is defined differently in every system, so the model learns three truths and trusts none.
- The process lives in the heads of the people who run it. When they move, it moves with them.
- The target operating model is a slide, not a system. Ownership is assumed rather than assigned.
- A migration moves records without a record of its own. Six months later nobody can say why a field was dropped or who agreed.

Each one is a missing layer underneath. The organisation built the top first, and the top fails first.

The spine: four layers, built from the bottom

Every engagement follows one order. A buyer may start at any layer. The work never skips one.

Layer	What it does	The rule	It fails when
1. Data	Holds the truth	Structure comes before any tool	The tool decides the model
2. Technology	Serves the truth	The platform follows the client's ecosystem	The platform choice drives the data model
3. Operating model	Keeps it true	Every field has an owner, and its classification is its access policy	Ownership lives in a slide
4. Outcome	Shows what the business gets	Good is written down before the work starts	Success is declared after the fact

Each layer rests on the one below it. AI is not a fifth layer on top. It is a capability inside Technology that reads the governed data and works within the operating model.

Principles

1. Define once. One definition of each customer, product and counterparty, owned, versioned and read by every consumer.
2. Keep the layer, replace the rest. Screens, reports, APIs, AI assistants and agents all read the governed layer. Each can be replaced without touching it.
3. Follow the client's platform. The method runs on what the client already has.
4. Rules first, AI where rules stop. Known logic is automated deterministically. AI handles what is ambiguous or too variable for rules.
5. Every statement keeps its source. Nothing is asserted without the document, record or line of code it came from.
6. Nothing moves until it clears. A separate check and a named person stand between a proposal and a change.
7. Write the outcome first. The measure and the date are agreed before the work, not found after it.

AI on evidence

A language model is good at reading and proposing. It is not a reliable judge of its own output. So the method separates the roles.

1. Read. The organisation's own material goes in: policies, procedures, contracts, tickets, transcripts, schemas and code.

2. Propose. AI proposes short, plain statements of fact, each tied to the exact passage it came from.
3. Check. A separate, independent check tests every statement against its source. Anything without a source is held as an open question, not a fact.
4. Clear. A named person in the business confirms or rejects what the check cannot settle. Conflicts between documents stay open until an owner decides.
5. Record. What was proposed, what checked it, which source it quotes, who cleared it and which version of the rules applied, all kept.

The statements that clear become registers the business already needs: a glossary with evidence, ownership, data classification, business rules, processes and roles. They land in the client's own catalogue and platform, versioned, and deletable by subject where personal data is involved.

The same discipline makes AI cheaper to trust. A reviewer checks the open questions, not every line.

Auditing transformations

A migration or a platform change is where meaning gets lost. Fields are mapped by people who were not there when the rules were written. Retention clauses are forgotten. Duplicates move twice.

The method puts a gate between the old estate and the new one. For every object in a wave, four questions are answered from the estate's own documents and code before it moves:

Question	Answered with
What is it?	Its classification and the policy clause behind it
Who signs for it?	The owner, with any disagreement left open until resolved
Which rules must it carry?	The business rules, each with the passage it came from
Should it move at all?	Move, archive or delete, with the retention clause

The gate connects nothing and moves nothing. It decides what may move and keeps the proof. When the auditor or the regulator asks why a field changed, the answer is a record, not a recollection.

How an engagement runs

Stage	Scope	Result
Review	Fixed scope, fixed duration, senior practitioners only	Current state, target state, the gap in each layer, a delivery plan
Pilot	One domain or one flow, on live infrastructure	One proven loop, measured against the definition of good
Scale	The same pattern on the next domains	A running system with owners in place, handed to the client's team

No pilot starts without a written scope, a named data owner, a record of every tool that touches the data, a test set the client agrees to judge, and an exit that names what is measured.

What the client keeps

The governed data layer on its own platform. The registers in its own catalogue. The operating model as runbooks, owners and gates. The full record of what changed and why. Nothing depends on Datapipdd to keep running.

Reported results

From delivered work between 2017 and 2026, as reported:

- 39 million dollars annual savings from an operations data programme.
- Data retrieval time down 70% after a platform move.
- More than 2 billion dollars of annual pass through cost made auditable end to end.

Contact

Start with a Review. pedro.duarte@datapipdd.com · datapipdd.com